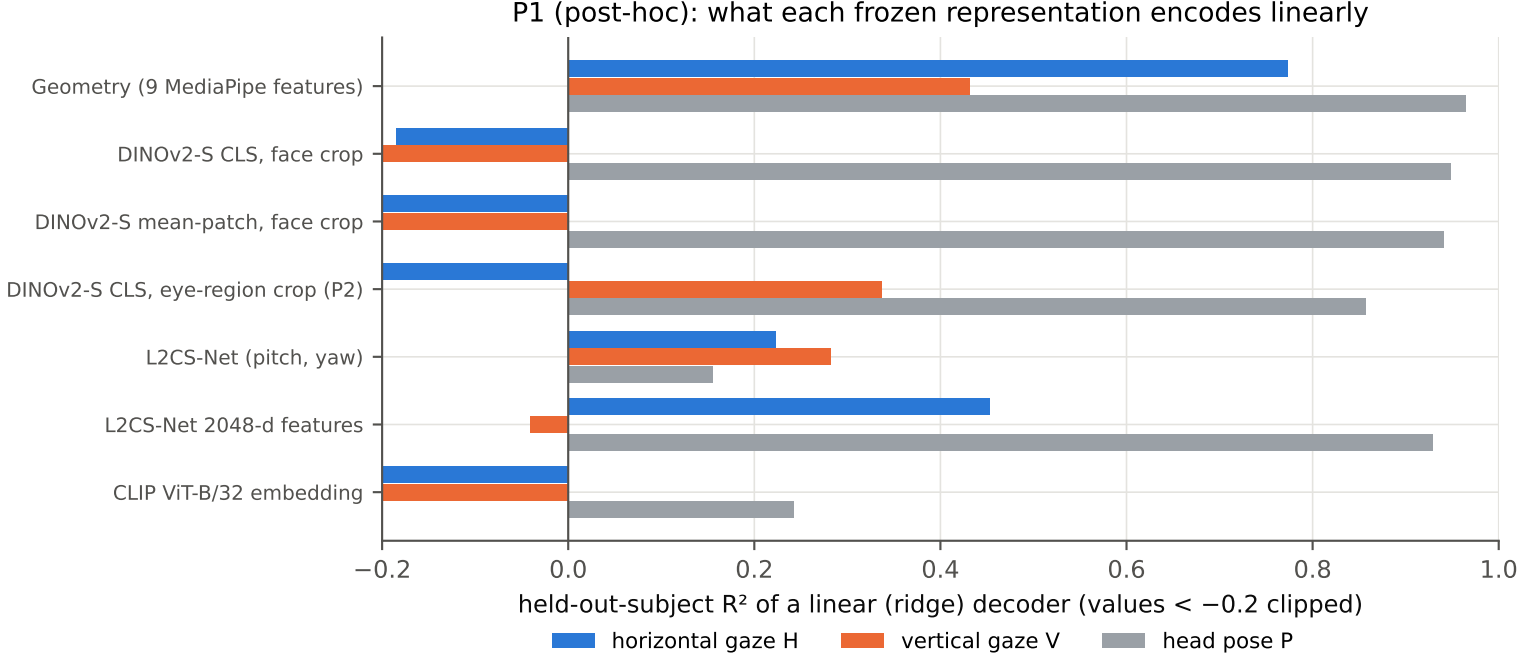
P1 (post-hoc): what each frozen representation encodes linearly
Geometry (9 MediaPipe features)
DINOv2-S CLS, face crop
DINOv2-S mean-patch, face crop
DINOv2-S CLS, eye-region crop (P2)
L2CS-Net (pitch, yaw)
L2CS-Net 2048-d features
CLIP ViT-B/32 embedding
−0.2
0.0
0.2
0.4
0.6
0.8
1.0
held-out-subject R² of a linear (ridge) decoder (values < −0.2 clipped)
horizontal gaze H
vertical gaze V
head pose P